

Importing libraries

```
In [168... import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

Importing data. Test data is not labelled or does not contain the churn column. Hence we will divide the train data into train and test.

```
In [169... df = pd.read_csv("train.csv")
test = pd.read_csv("test.csv")
df.shape
```

```
Out[169... (4250, 20)
```

Exploratory Data Analysis

```
In [170... df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4250 entries, 0 to 4249
Data columns (total 20 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   state                                  4250 non-null   object
 1   account_length                        4250 non-null   int64
 2   area_code                             4250 non-null   object
 3   international_plan                    4250 non-null   object
 4   voice_mail_plan                       4250 non-null   object
 5   number_vmail_messages                 4250 non-null   int64
 6   total_day_minutes                     4250 non-null   float64
 7   total_day_calls                       4250 non-null   int64
 8   total_day_charge                      4250 non-null   float64
 9   total_eve_minutes                     4250 non-null   float64
10   total_eve_calls                       4250 non-null   int64
11   total_eve_charge                      4250 non-null   float64
12   total_night_minutes                   4250 non-null   float64
13   total_night_calls                     4250 non-null   int64
14   total_night_charge                    4250 non-null   float64
15   total_intl_minutes                     4250 non-null   float64
16   total_intl_calls                       4250 non-null   int64
17   total_intl_charge                     4250 non-null   float64
18   number_customer_service_calls         4250 non-null   int64
19   churn                                 4250 non-null   object
dtypes: float64(8), int64(7), object(5)
memory usage: 664.2+ KB
```

Checking for null values

```
In [171... df.isnull().sum()
```

```
Out[171... state                                0
account_length                          0
area_code                               0
international_plan                       0
voice_mail_plan                         0
number_vmail_messages                    0
total_day_minutes                        0
total_day_calls                          0
total_day_charge                         0
total_eve_minutes                        0
total_eve_calls                          0
total_eve_charge                         0
total_night_minutes                      0
total_night_calls                        0
```

```

total_night_charge      0
total_intl_minutes      0
total_intl_calls        0
total_intl_charge       0
number_customer_service_calls  0
churn                   0
dtype: int64

```

Checking for any redundant or abnormal data in categorical variables. Finding - Our target variable "churn" is highly imbalanced. This could create a bad prediction on class yes.

```

In [172... for x in df.columns:
            if df[x].dtype == "object":
                print(x)
                print(df[x].value_counts(), "\n")

```

```

state
WV      139
MN      108
ID       96
AL       95
VA       94
OR       93
TX       92
UT       91
NY       90
NJ       89
OH       88
WY       87
WI       86
MA       85
ME       84
CT       83
RI       82
MI       81
KS       80
MD       79
VT       78
KY       77
IN       76
NV       75
MS       74
DE       73
WA       72
MT       71
MO       70
NC       69
CO       68
IL       67
TN       66
NH       65
OK       64
NM       63
HI       62
AZ       61
FL       60
SD       59
NE       58
SC       57
DC       56
AR       55
LA       54
ND       53
PA       52
GA       51
IA       50
AK       49
CA       48
Name: state, dtype: int64

```

```

area_code
area_code_415    2108
area_code_408    1086
area_code_510    1056
Name: area_code, dtype: int64

international_plan
no      3854
yes     396
Name: international_plan, dtype: int64

voice_mail_plan
no      3138
yes     1112
Name: voice_mail_plan, dtype: int64

churn
no      3652
yes     598
Name: churn, dtype: int64

```

Feature engineering and reduction: Here we have combined the calls, charges and minutes for the entire day.

```

In [173... df["total_calls"] = df["total_day_calls"] + df["total_eve_calls"] + df["total_night_
df["total_charge"] = df["total_day_charge"] + df["total_eve_charge"] + df["total_nig
df["total_minutes"] = df["total_day_minutes"] + df["total_eve_minutes"] + df["total_

df = df.drop(['total_day_calls', 'total_eve_calls', 'total_night_calls', 'total_day_cha

```

Separating categorical and numerical variables and storing in a list.

```

In [174... cat_var=[]
num_var=[]

for x in df.columns:
    if df[x].dtype == "int64" or df[x].dtype == "float64":
        num_var.append(x)
    else:
        cat_var.append(x)

```


Visualizing numerical data for their distributions. number_vmail_messages has most of the data as 0. Hence dropping this column

```

In [175... j=1
plt.figure(figsize=(35,35))
for x in num_var:
    plt.subplot(10,10,j)
    plt.hist(df[x])
    plt.xlabel(x)
    j+=1

#number_vmail_messages has most of entries as 0

```



```

In [176... df = df.drop(["number_vmail_messages"], axis = 1)

```

Plotting 1st correlation matrix to identify multicollinearity. Found that charge and minutes\ are

multi collinear and hence dropping the minutes variable.

```
In [177... corr_df=df.iloc[:,:].corr(method="pearson")
# print(corr_df)
plt.figure(figsize = (30,10)) #Change the size of heatmap
sns.heatmap(corr_df,vmax=1.0,vmin=-1.0,annot=True,linewidths=.5)
```

Out[177... <AxesSubplot:>



```
In [178... df = df.drop(["total_minutes","total_intl_minutes"],axis=1)
```

Label encoding all categorical variables

```
In [179... from sklearn import preprocessing
```

```
le = preprocessing.LabelEncoder()
for x in cat_var:
    df[x] = le.fit_transform(df[x])
    le_mapping = dict(zip(le.classes_,le.transform(le.classes_)))
    print(x,le_mapping)
```

```
state {'AK': 0, 'AL': 1, 'AR': 2, 'AZ': 3, 'CA': 4, 'CO': 5, 'CT': 6, 'DC': 7, 'DE': 8, 'FL': 9, 'GA': 10, 'HI': 11, 'IA': 12, 'ID': 13, 'IL': 14, 'IN': 15, 'KS': 16, 'KY': 17, 'LA': 18, 'MA': 19, 'MD': 20, 'ME': 21, 'MI': 22, 'MN': 23, 'MO': 24, 'MS': 25, 'MT': 26, 'NC': 27, 'ND': 28, 'NE': 29, 'NH': 30, 'NJ': 31, 'NM': 32, 'NV': 33, 'NY': 34, 'OH': 35, 'OK': 36, 'OR': 37, 'PA': 38, 'RI': 39, 'SC': 40, 'SD': 41, 'TN': 42, 'TX': 43, 'UT': 44, 'VA': 45, 'VT': 46, 'WA': 47, 'WI': 48, 'WV': 49, 'WY': 50}
area_code {'area_code_408': 0, 'area_code_415': 1, 'area_code_510': 2}
international_plan {'no': 0, 'yes': 1}
voice_mail_plan {'no': 0, 'yes': 1}
churn {'no': 0, 'yes': 1}
```

Plotting 2nd correlation matrix with all categorical variables to identify relation with target variable.\ Findings - International_plan, number_customer_service_calls and total_charge have a good correlation with the target variable.

```
In [180... corr_df=df.iloc[:,:].corr(method="pearson")
# print(corr_df)
plt.figure(figsize = (30,10)) #Change the size of heatmap
sns.heatmap(corr_df,vmax=1.0,vmin=-1.0,annot=True,linewidths=.5)
```

Out[180... <AxesSubplot:>



Separating the churn variable with rest of the variables.

```
In [181... X = df.loc[:,df.columns != 'churn']
Y = df.loc[:,df.columns == 'churn']
```

Performing train test split

```
In [182... from sklearn.model_selection import train_test_split

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size =0.3, random_sta
```

Logistic Regression

```
In [183... from sklearn.linear_model import LogisticRegression

model = LogisticRegression()
model.fit(X_train,Y_train)
model.score(X_train,Y_train)
```

C:\Users\PARSHVA\anaconda3\lib\site-packages\sklearn\utils\validation.py:72: DataCon
versionWarning: A column-vector y was passed when a 1d array was expected. Please ch
ange the shape of y to (n_samples,), for example using ravel().
return f(**kwargs)

C:\Users\PARSHVA\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:762:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

```
Out[183... 0.865546218487395
```

```
In [184... Y_pred = model.predict(X_test)
```

```
In [185... from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

print(confusion_matrix(Y_test,Y_pred))
print()
print(accuracy_score(Y_test,Y_pred))
print(classification_report(Y_test,Y_pred))
```

```
[[1074  21]
 [ 147  33]]
```

```
0.8682352941176471
```

```
precision    recall  f1-score   support
```

```
0          0.88          0.98          0.93         1095
```

1	0.61	0.18	0.28	180
accuracy			0.87	1275
macro avg	0.75	0.58	0.60	1275
weighted avg	0.84	0.87	0.84	1275

Random Forest

```
In [186... from sklearn.ensemble import RandomForestClassifier
clf = RandomForestClassifier()
clf.fit(X_train,Y_train )
clf.score(X_train, Y_train)
```

<ipython-input-186-5e7d161455bb>:3: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
clf.fit(X_train,Y_train )
```

Out[186... 1.0

```
In [187... Y_pred = clf.predict(X_test)
print(confusion_matrix(Y_test,Y_pred))
print()
print(accuracy_score(Y_test,Y_pred))
print(classification_report(Y_test,Y_pred))
```

```
[[1095    0]
 [  35  145]]
```

0.9725490196078431

	precision	recall	f1-score	support
0	0.97	1.00	0.98	1095
1	1.00	0.81	0.89	180
accuracy			0.97	1275
macro avg	0.98	0.90	0.94	1275
weighted avg	0.97	0.97	0.97	1275

The most important features for our model were: total_charge international_plan
number_customer_service_calls

Visualizations

Importing libraries

```
In [4]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

import warnings
warnings.filterwarnings("ignore")

pd.set_option('max_columns',100)
pd.set_option('max_rows',900)

pd.set_option('max_colwidth',200)

df = pd.read_csv("train.csv")
```

Separating numerical and categorical columns

```
In [6]: numerical = df.select_dtypes(include = "number").columns
cat = df.select_dtypes(include = "object").columns
```

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4250 entries, 0 to 4249
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                4250 non-null   object
1   account_length                       4250 non-null   int64
2   area_code                            4250 non-null   object
3   international_plan                   4250 non-null   object
4   voice_mail_plan                     4250 non-null   object
5   number_vmail_messages               4250 non-null   int64
6   total_day_minutes                   4250 non-null   float64
7   total_day_calls                     4250 non-null   int64
8   total_day_charge                    4250 non-null   float64
9   total_eve_minutes                   4250 non-null   float64
10  total_eve_calls                     4250 non-null   int64
11  total_eve_charge                    4250 non-null   float64
12  total_night_minutes                 4250 non-null   float64
13  total_night_calls                   4250 non-null   int64
14  total_night_charge                  4250 non-null   float64
15  total_intl_minutes                  4250 non-null   float64
16  total_intl_calls                    4250 non-null   int64
17  total_intl_charge                   4250 non-null   float64
18  number_customer_service_calls       4250 non-null   int64
19  churn                              4250 non-null   object
dtypes: float64(8), int64(7), object(5)
memory usage: 664.2+ KB
```

```
In [8]: df1 = df.copy()
```

Encoding categorical data so that model can understand it. Here all the text gets a number as shown below.

```
In [10]: from sklearn import preprocessing
```

```
le = preprocessing.LabelEncoder()
for x in cat:
    df1[x] = le.fit_transform(df1[x])
    le_mapping = dict(zip(le.classes_, le.transform(le.classes_)))
    print(x, le_mapping)
```

```
state {'AK': 0, 'AL': 1, 'AR': 2, 'AZ': 3, 'CA': 4, 'CO': 5, 'CT': 6, 'DC': 7, 'DE': 8, 'FL': 9, 'GA': 10, 'HI': 11, 'IA': 12, 'ID': 13, 'IL': 14, 'IN': 15, 'KS': 16, 'KY': 17, 'LA': 18, 'MA': 19, 'MD': 20, 'ME': 21, 'MI': 22, 'MN': 23, 'MO': 24, 'MS': 25, 'MT': 26, 'NC': 27, 'ND': 28, 'NE': 29, 'NH': 30, 'NJ': 31, 'NM': 32, 'NV': 33, 'NY': 34, 'OH': 35, 'OK': 36, 'OR': 37, 'PA': 38, 'RI': 39, 'SC': 40, 'SD': 41, 'TN': 42, 'TX': 43, 'UT': 44, 'VA': 45, 'VT': 46, 'WA': 47, 'WI': 48, 'WV': 49, 'WY': 50}
area_code {'area_code_408': 0, 'area_code_415': 1, 'area_code_510': 2}
international_plan {'no': 0, 'yes': 1}
voice_mail_plan {'no': 0, 'yes': 1}
churn {'no': 0, 'yes': 1}
```

Percentage churning based on total number of service calls Finding - When number of service calls exceed 4 the chances of churn is more than 44%.

```
In [14]: sc=[]
rate=[]
for x in df["number_customer_service_calls"].unique():
    sc.append(x)
    rate.append(df[df["number_customer_service_calls"]== x]['churn'].mean()*100)


pd.DataFrame(list(zip(sc,rate)), columns=['service calls','churn rate']).sort_values
```

```
Out[14]:
```

	service calls	churn rate
7	9	100.000000
8	6	67.857143
5	5	60.493827
6	7	53.846154
9	8	50.000000
4	4	44.019139
3	3	11.290323
1	0	10.948081
0	1	10.892388
2	2	10.770855

%people churning with an international plan Findings - If a person has an international plan then he has a probability of 42.2% of churning.

```
In [16]: plt.figure(figsize=(8,5))
ax = sns.countplot(data=df1[df1["international_plan"] == 1],x="churn")
for p in ax.patches:
    percentage = '{:.1f}%'.format(100 * p.get_height()/len(df1[df1["international_pl
x = p.get_x() + 0.45
y = p.get_height()
ax.annotate(percentage, (x, y),ha='center')
plt.show()
```

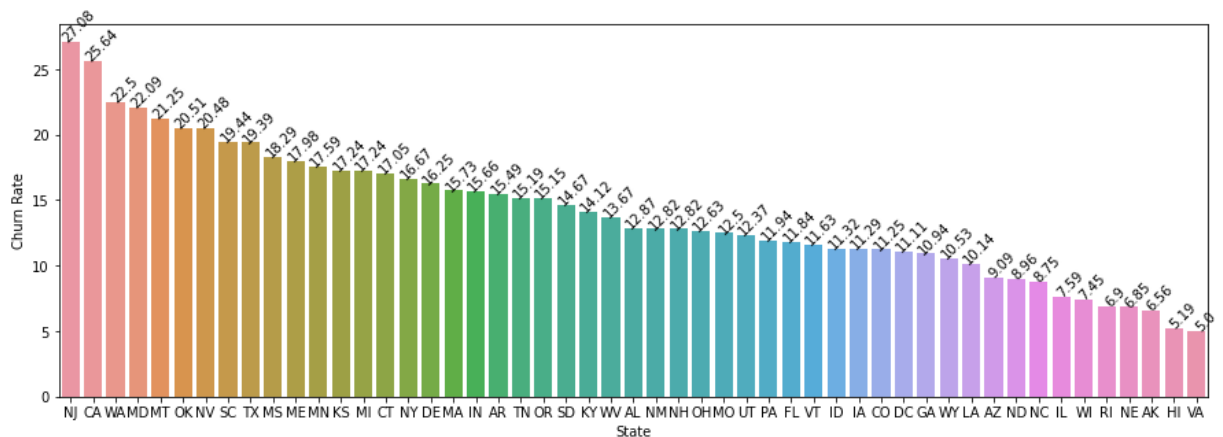



State wise highest churn rate

```
In [18]: state=[]
percentage_churn=[]
for x in df['state'].unique():
    state.append(x)
    percentage_churn.append((len(df[df['state']==x][df['churn']=='yes'])/len(df[df['state']==x]))

state_churnrate = pd.DataFrame(list(zip(state,percentage_churn)),columns=['State','Churn Rate'])

plt.figure(figsize=(15,5))
ax=sns.barplot(data=state_churnrate,x="State",y="Churn Rate")
for p in ax.patches:
    x = p.get_x()+ p.get_width()
    y = p.get_height()
    ax.annotate(round(p.get_height(),2), (x, y),ha='center').set_rotation(45)
plt.show()
```



```
In [19]: df1["total_calls"] = df1["total_day_calls"] + df1["total_eve_calls"] + df1["total_ni
df1["total_charge"] = df1["total_day_charge"] + df1["total_eve_charge"] + df1["total
df1["total_minutes"] = df1["total_day_minutes"] + df1["total_eve_minutes"] + df1["to

#df1 = df1.drop(['total_day_calls', 'total_eve_calls', 'total_night_calls', 'total_day_
```


Number of churned customers vs total charges. Most of the customers churn when they cross a spend between 70 to 80 dollars

```
In [42]: plt.figure(figsize=(16,8))
```

```
plt.hist(data=df1[df1["churn"]==1],x="total_charge",linewidth=1.0,edgecolor='black',  
plt.xlabel("total_charge"))
```

#Majority of churning happens when customer crosses 70 to 80 dollars of talktime cha

Out[42]: Text(0.5, 0, 'total_charge')



In []: